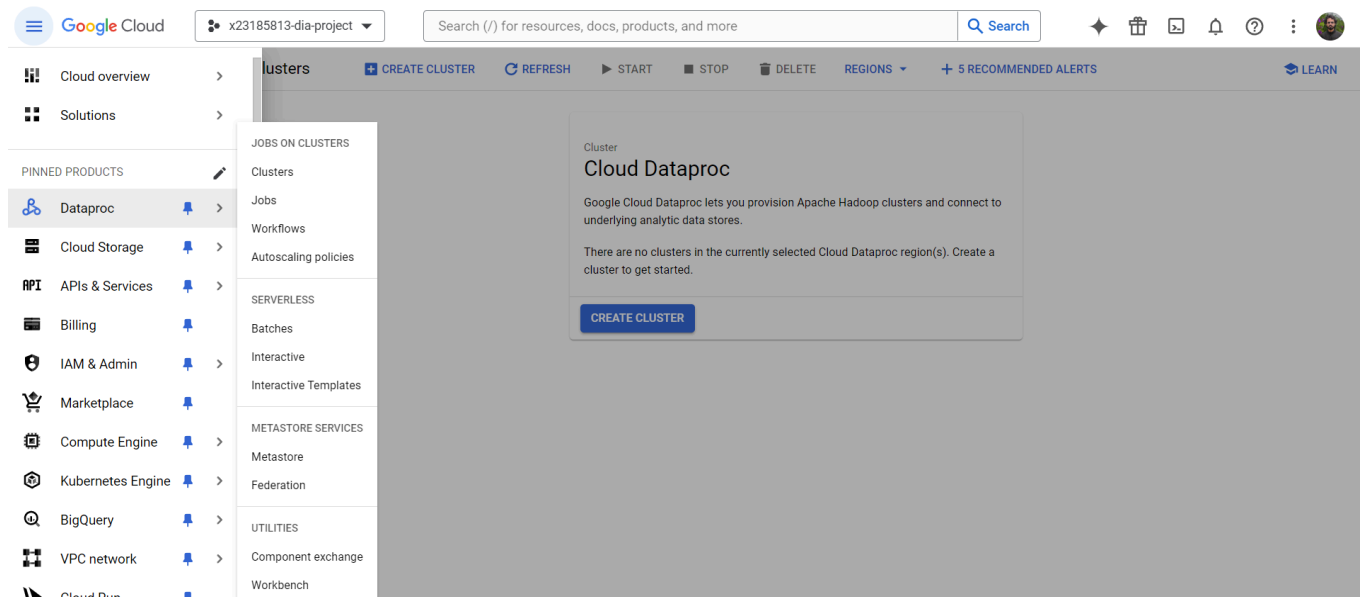


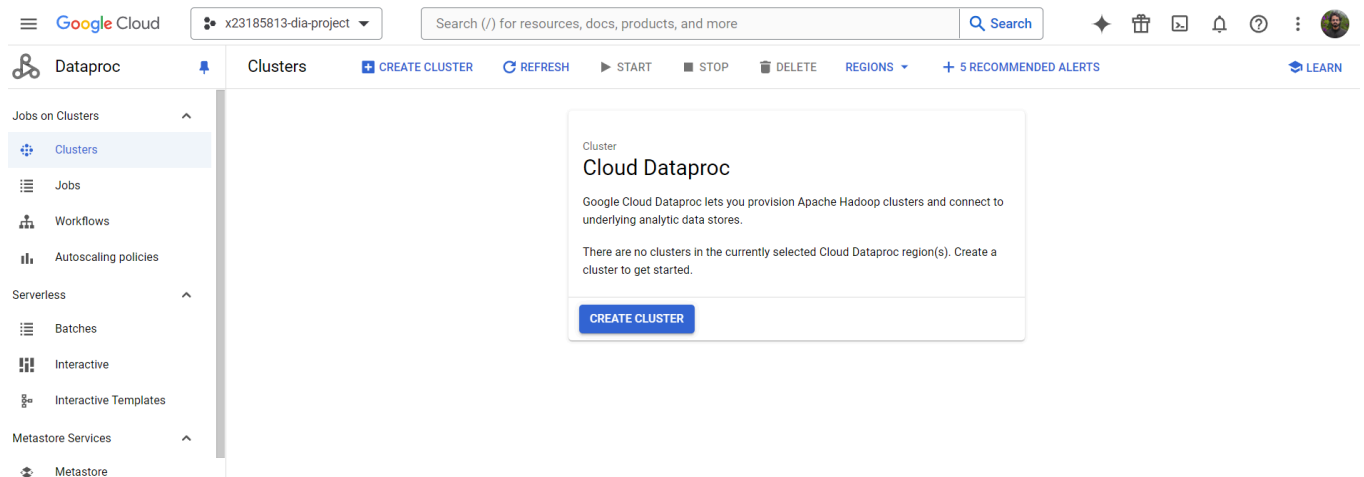
# Analysis of Motor Vehicle Collisions in New York City with PySpark and MapReduce

This guide will help setup the dataproc cluster and apache spark used for this project. We use Google Cloud for this, so as prerequisite you need a google account and a billing account associated with it to access the google cloud console

1. Log in to the google cloud console and select dataproc from the side panel or use the search bar



1. The dataproc interface will look like this proceed with **CREATE CLUSTER**



3. Select Cluster on Compute engine this will use the google compute engine to create the virtual machines used for master and worker nodes

### Create Dataproc cluster

Select the infrastructure service that you want to use.

**Cluster on Compute Engine**  
Create the cluster on Compute Engine.

CREATE

**Cluster on GKE**  
Create the cluster on Google Kubernetes Engine (GKE).

CREATE

CANCEL

4. In the following steps select the settings for the cluster.

select a name for the cluster and select the cluster type we use 'Standard'

## Name

Cluster Name \*  
cluster-dia-x23185813 ?

## Location

Region \*  
europe-west1 ▼ ?

Zone \*  
Any ▼ ?

## Cluster type

- ☒ Standard (1 master, N workers)
- ☐ Single Node (1 master, 0 workers)  
Provides one node that acts as both master and worker. Good for proof-of-concept or small-scale processing
- ☐ High Availability (3 masters, N workers)  
Hadoop High Availability mode provides uninterrupted YARN and HDFS operations despite single-node failures or reboots

## Versioning

Use a custom image to load pre-installed packages. [Learn more](#)

### Image Type and Version

2.2-debian12

### Release Date

First released on 12/08/20

CHANGE

Choose the subnetwork as default

## Network Configuration



Establishes connectivity for the VM instances in this cluster.

☒ Networks in this project

[Learn more](#)

☐ Networks shared from host project: ""

Choose a shared VPC network from project that is different from this cluster's project.

[Learn more](#)

Primary network

default



Subnetwork

default



Network tags

Network tags are text attributes you can add to make firewall rules and routes applicable to specific VM instances.

## Dataproc Metastore

Configure Dataproc to use Dataproc Metastore as its Hive metastore. [Learn more](#)

Selected project

x23185813-dia-project

[BROWSE](#)

Metastore service

None



We recommend this option to persist table metadata when a cluster is shut down, for a metastore shared by different clusters, or for metadata operability across GCP products.

Check the Component gateway and select Jupyter Notebook

## Components

### Component Gateway

☒ Enable component gateway

Provides access to the web interfaces of default and selected optional components on the cluster. [Learn more](#)

### Optional components

Select one or multiple components. [Learn more](#)

☐ Anaconda ?

☐ Hive WebHCat ?

☒ Jupyter Notebook ?

☐ Zeppelin Notebook ?

☐ Trino ?

☐ ZooKeeper ?

☐ Ranger ?

☐ Flink ?

☐ Docker ?

☐ Solr ?

☐ Hudi ?

• Set up cluster  
Begin by providing basic information.

• **Configure nodes (optional)**  
Change node compute and storage capabilities.

• Customize cluster (optional)  
Add cluster properties, features, and actions.

• Manage security (optional)  
Change access, encryption, and security settings.

CREATE

CANCEL

EQUIVALENT COMMAND LINE ▾

In the Configure node select the node configurations

we will use E2 instances

1 Master node and 2 Worker nodes

## Manager node



Contains the YARN Resource Manager, HDFS NameNode, and all job drivers.

✓ General purpose

Compute optimized

Memory optimized

GPUs

Machine types for common workloads, optimized for cost and flexibility

Series

E2 ▾

CPU platform selection based on availability

Machine type

e2-standard-2 (2 vCPU, 1 core, 8 GB memory) ▾



vCPU  
2

Memory  
8 GB

### ✓ CPU PLATFORM AND GPU

Primary disk size \*

200 GB ?

Primary disk type \*

Balanced Persistent Disk ▾ ?

Number of local SSDs ▾ x 375GB ?

Local SSD Interface ▾ ?

## Worker nodes



Each contains a YARN NodeManager and a HDFS DataNode. HDFS replication factor is 2.

☒ General purpose

Compute optimized

Memory optimized

GPUs

Machine types for common workloads, optimized for cost and flexibility

Series

E2

CPU platform selection based on availability

Machine type

e2-standard-2 (2 vCPU, 1 core, 8 GB memory)



vCPU

2

Memory

8 GB

### ✓ CPU PLATFORM AND GPU

Number of worker nodes \*

2

?

Primary disk size \*

100

GB

?

Primary disk type \*

Balanced Persistent Disk

?

Number of local SSDs

x 375GB

?

Local SSD Interface

?

## Secondary worker nodes



Each contains a YARN NodeManager. HDFS does not run on secondary worker nodes. Secondary worker VMs are preemptible by default. Spot and preemptible VMs costs less, but can be terminated at any time due to system demands. [Learn more](#)

## Sole-tenancy

Enable to create this cluster on sole-tenant nodes. This grants exclusive access to a physical Compute Engine server that is dedicated to hosting only your project's VMs. If you are creating a cluster with an autoscaling policy, it is recommended that the node group you select also uses an autoscaling policy. [Learn more](#)

☐ Enable

## Shielded VM

Turn on all settings for the most secure configuration. [Learn more](#)


- ☐ Turn on Secure Boot ?
- ☐ Turn on vTPM ?
- ☐ Turn on Integrity Monitoring ?

## Total YARN usage

YARN cores ?	YARN memory ?
4	12.8 GB

Uncheck the internal IP only option

## Internal IP only

☐ Configure all instances to have only internal IP addresses. [Learn more](#) 

## Labels

A list of key:value pairs to attach to the cluster for tracking.

[+ ADD LABELS](#)

## Cluster properties

Use cluster properties to add or modify configuration files when creating a cluster.

[+ ADD PROPERTIES](#)

## Initialization actions

Use initialization actions to customize settings, install applications, or make other modifications to your cluster. Select scripts or executables that Cloud Dataproc will run when provisioning your cluster.

[+ ADD INITIALIZATION ACTION](#)


## Custom cluster metadata

Add custom metadata to cluster instances. [Learn more](#) 


[+ ADD METADATA](#)

Rest settings keep it as default

## Project access

☐ Enables the cloud-platform scope for this cluster [Learn more](#) 

## Encryption

Encrypt cluster persistent disk data and optionally job argument data. [Learn more](#) 

☒ Google-managed encryption key


Keys owned by Google

☐ Cloud KMS key


Keys owned by customers

☐ Encrypt job argument data in addition to cluster persistent disk data.

☐ Enable confidential computing


Confidential Computing on clusters can only be enabled if all nodes on the cluster use the N2D machine type. [Learn more](#) 

## Personal Cluster Authentication

Enable Dataproc Personal Cluster Authentication to allow interactive workloads on the cluster to securely run as your end user identity. [Learn more](#) 

☐ Enable

## Secure Multi Tenancy

Enable Dataproc Service Account Based Secure Multi-tenancy to share a cluster with multiple users. Make sure the VM service account for the cluster has the proper permissions to impersonate all mapped service accounts for users. [Learn more](#) 

☐ Enable

5. Create the cluster using **CREATE**

- **Set up cluster**  
Begin by providing basic information.
- **Configure nodes (optional)**  
Change node compute and storage capabilities.
- **Customize cluster (optional)**  
Add cluster properties, features, and actions.
- **Manage security (optional)**  
Change access, encryption, and security settings.

CREATE

CANCEL

EQUIVALENT COMMAND LINE ▾

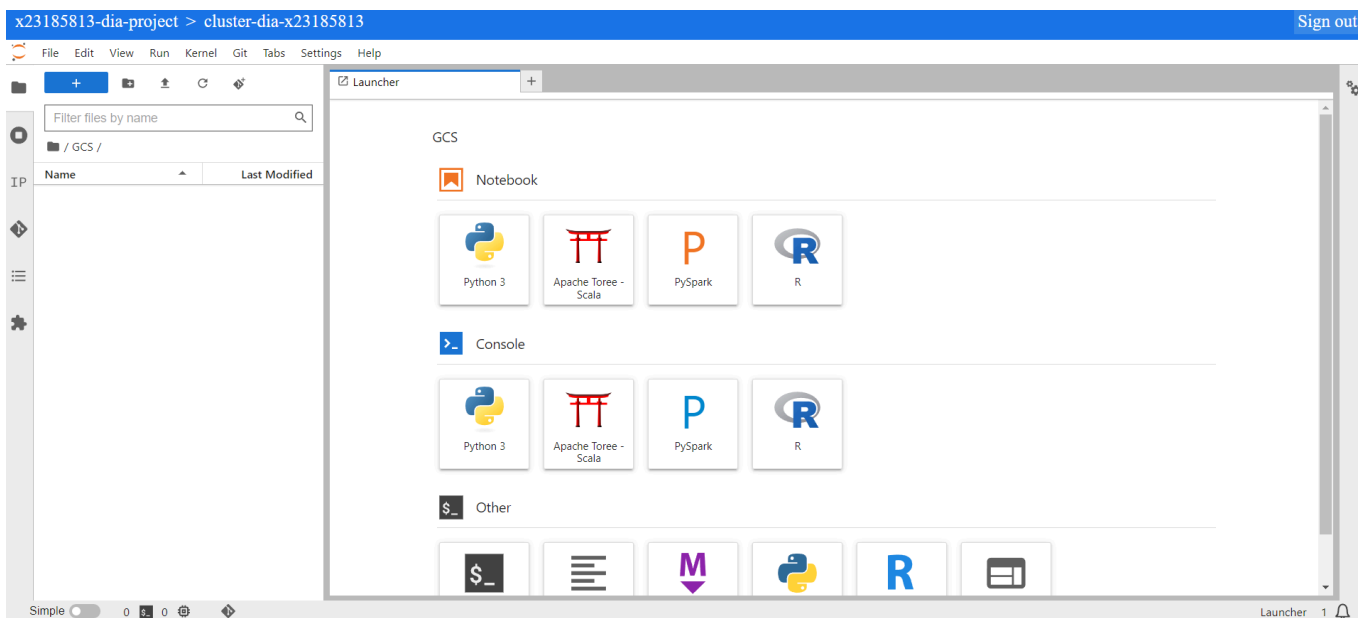
The dataproc will start provisioning the resources and will install all the required softwares.

The screenshot shows the Google Cloud Dataproc console for a cluster named 'cluster-dia-x23185813'. The cluster is in the 'Provisioning' status. A warning message states: 'For PD-Standard without local SSDs, we strongly recommend provisioning 1TB or larger to ensure consistently high I/O performance. See https://cloud.google.com/compute/docs/disks/performance for information on disk I/O performance.' The 'VM INSTANCES' tab is selected, showing three instances: one Master node (cluster-dia-x23185813-m) and two Worker nodes (cluster-dia-x23185813-w-0 and cluster-dia-x23185813-w-1), all using 'e2-standard-2' machine type. A toast notification at the bottom says 'Request to create cluster cluster-dia-x23185813 submitted'.

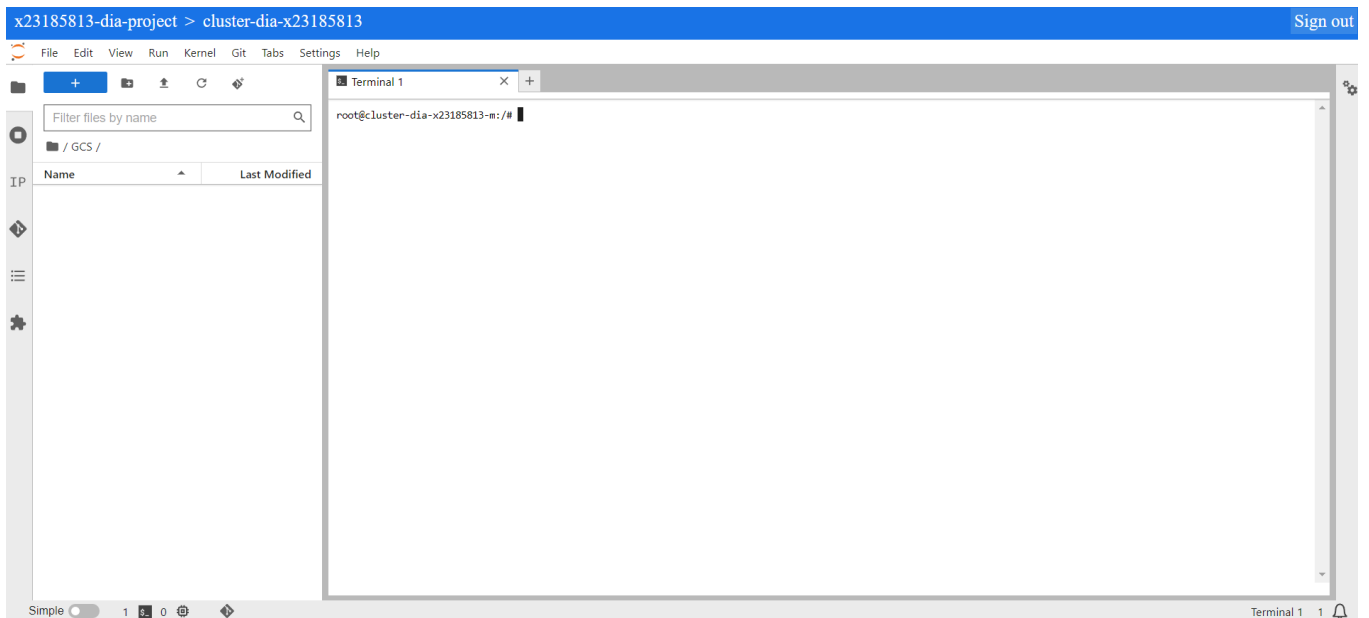
6. Once the Cluster is up and running go to **WEB INTERFACES** select jupyterlab to open the jupyter interface

The screenshot shows the same Dataproc cluster, but now its status is 'Running'. The 'WEB INTERFACES' tab is selected, displaying a list of available interfaces: 'SSH tunnel', 'Component gateway', 'YARN ResourceManager', 'MapReduce Job History', 'Spark History Server', 'HDFS NameNode', 'YARN Application Timeline', 'Tez', 'Jupyter', and 'JupyterLab'. The 'JupyterLab' option is highlighted at the bottom of the list.

7. The jupyterLab interface will look like this, In launcher select the console option



This will open a terminal window with root user



8. We will execute the following commands to setup the files from the git repository

```
su dataproc
```

```
cd ~
```

```
git clone https://github.com/ankithbjoseph/dia-2024-project.git
```

```
cd dia-2024-project/
```



The screenshot shows a JupyterLab interface with a terminal window open. The terminal output shows the following commands and their results:

```
root@cluster-dia-x23185813-m:/# su dataproc
dataproc@cluster-dia-x23185813-m:/$ ls
bin boot copyright dev etc hadoop home lib lib64 lost+found media mnt opt proc root run sbin srv sys usr var
dataproc@cluster-dia-x23185813-m:/$ cd ~
dataproc@cluster-dia-x23185813-m:/$ ls
dataproc@cluster-dia-x23185813-m:/$ git clone https://github.com/ankithjoseph/dia-2024-project.git
Cloning into 'dia-2024-project'...
remote: Enumerating objects: 10, done.
remote: Counting objects: 100% (10/10), done.
remote: Compressing objects: 100% (9/9), done.
remote: Total 10 (delta 2), reused 5 (delta 0), pack-reused 0
Receiving objects: 100% (10/10), 211.45 KiB | 6.04 MiB/s, done.
Resolving deltas: 100% (2/2), done.
dataproc@cluster-dia-x23185813-m:/$ cd dia-2024-project/
dataproc@cluster-dia-x23185813-m:~/dia-2024-project$ ls
README.md x23185813.ipynb
dataproc@cluster-dia-x23185813-m:~/dia-2024-project$
```

9. This will download the required files , In the file explorer navigate to the folder `home/dataproc/dia-2024-project/`

10. Open the notebook `x23185813.ipynb`

The screenshot shows a JupyterLab interface with a notebook open. The notebook cell contains the following code:

```
[1]: CRASHES_URL = "https://data.cityofnewyork.us/api/views/h9gi-nx95/rows.csv?accessType=DOWNLOAD"
VEHICLES_URL = "https://data.cityofnewyork.us/api/views/bm4k-52h4/rows.csv?accessType=DOWNLOAD"
BUCKET_NAME = "bucket-dia-x23185813"
CRASHES_FILE_NAME = "Crashes.csv"
VEHICLES_FILE_NAME = "Vehicles.csv"
PATH = "notebooks/jupyter/dataset"

# Check if the Crashes file exists in the bucket and download if it does not
crashes_exists = !gsutil ls gs://{BUCKET_NAME}/{PATH}/{CRASHES_FILE_NAME}
if 'CommandException: One or more URLs matched no objects.' in crashes_exists:
    !wget {CRASHES_URL} -O /tmp/{CRASHES_FILE_NAME}
    !gsutil cp /tmp/{CRASHES_FILE_NAME} gs://{BUCKET_NAME}/{PATH}/{CRASHES_FILE_NAME}
    !rm /tmp/{CRASHES_FILE_NAME}
else:
    print("Crashes dataset already exists in GCS. Skipping download.")

# Check if the Vehicles file exists in the bucket and download if it does not
vehicles_exists = !gsutil ls gs://{BUCKET_NAME}/{PATH}/{VEHICLES_FILE_NAME}
if 'CommandException: One or more URLs matched no objects.' in vehicles_exists:
    !wget {VEHICLES_URL} -O /tmp/{VEHICLES_FILE_NAME}
    !gsutil cp /tmp/{VEHICLES_FILE_NAME} gs://{BUCKET_NAME}/{PATH}/{VEHICLES_FILE_NAME}
    !rm /tmp/{VEHICLES_FILE_NAME}
else:
    print("Vehicles dataset already exists in GCS. Skipping download.")

Crashes dataset already exists in GCS. Skipping download.
Vehicles dataset already exists in GCS. Skipping download.

Load libraries and define functions
```

11. Run All cells

x23185813-dia-project > cluster-dia-x23185813

Sign out

File Edit View Run Kernel Git Tabs Settings Help

+ Filter files by name

/ ... / dataproc

Name

README.md

x23185813.ipynb

Run Selected Cells

Run Selected Cells and Insert Below

Run Selected Cells and Do not Advance

Run Selected Text or Current Line in Console

Run All Above Selected Cell

Run Selected Cell and All Below

Render All Markdown Cells

Run All Cells

Restart Kernel and Run All Cells...

Shift+Enter

Alt+Enter

Ctrl+Enter

Code

git

PySpark

```
"https://data.cityofnewyork.us/api/views/h9gi-nx95/rows.csv?accessType=DOWNLOAD"
"https://data.cityofnewyork.us/api/views/bm4k-52h4/rows.csv?accessType=DOWNLOAD"
bucket-dia-x23185813"
NAME = "Crashes.csv"
NAME = "Vehicles.csv"
kks/jupyter/dataset"

Crashes file exists in the bucket and download if it does not
= lgsutil ls gs://{BUCKET_NAME}/{PATH}/{CRASHES_FILE_NAME}
ception: One or more URLs matched no objects.' in crashes_exists:
lwget {CRASHES_URL} -O /tmp/{CRASHES_FILE_NAME}
lgsutil cp /tmp/{CRASHES_FILE_NAME} gs://{BUCKET_NAME}/{PATH}/{
lrm /tmp/{CRASHES_FILE_NAME}

else:
    print("Crashes dataset already exists in GCS. Skipping download.")

# Check if the Vehicles file exists in the bucket and download if it does not
vehicles_exists = lgsutil ls gs://{BUCKET_NAME}/{PATH}/{VEHICLES_FILE_NAME}
if 'CommandException: One or more URLs matched no objects.' in vehicles_exists:
    lwget {VEHICLES_URL} -O /tmp/{VEHICLES_FILE_NAME}
    lgsutil cp /tmp/{VEHICLES_FILE_NAME} gs://{BUCKET_NAME}/{PATH}/{
    lrm /tmp/{VEHICLES_FILE_NAME}

else:
    print("Vehicles dataset already exists in GCS. Skipping download.")

Crashes dataset already exists in GCS. Skipping download.
Vehicles dataset already exists in GCS. Skipping download.
```

Load libraries and define functions

Simple 1 1 1 1 PySpark | Idle

Mode: Command Ln 1, Col 1 x23185813.ipynb 1